



Software Architecture Laboratory

Laboratory No. 4 REST API – Blueprints

Blueprints Java 21 / Spring Boot 3.3.x

Students:

Andersson David Sánchez Méndez

Cristian Santiago Pedraza Rodríguez

Elizabeth Correa Suárez

Juan Sebastian Ortega Muñoz

Instructor: Professor Javier Ivan Toquica Barrera

Course: Software Architecture (ARSW)

Institution: Escuela Colombiana de Ingeniería Julio Garavito

February 2026

Contents

1 Objective	3
2 Application Architecture	3
2.1 Package Structure	3
2.2 Request / Response Flow	3
3 Part 1: Code Base Familiarization	3
3.1 Model Layer	3
3.2 Persistence Layer	3
3.3 Service Layer	3
3.4 Controller Layer	4
4 Part 2: Migration to PostgreSQL	4
4.1 Design Decisions	4
4.2 JPA Entity Model	4
4.3 Auto-Generated Relational Model	4
4.4 Domain ↔ Entity Mapping	4
4.5 Docker Setup	4
4.6 Database Verification	4
5 Part 3: REST API Best Practices	5
5.1 Versioned Base Path	5
5.2 Uniform Response Wrapper	5
5.3 HTTP Status Code Mapping	5
5.4 Error Response Format	5
6 Part 4: OpenAPI / Swagger	5
6.1 Configuration	5
6.2 Endpoint Annotations	5
6.3 Swagger UI Evidence	6
7 Part 5: Blueprints Filters	6
7.1 Design Pattern	6
7.2 Available Filters	6
7.3 RedundancyFilter	6
7.4 UndersamplingFilter	6
7.5 Profile-Based Activation	6
7.6 Filter Evidence	7
8 Testing Strategy	7
8.1 Test Suite Overview	7
8.2 Notable Test Techniques	7
8.3 SonarCloud Coverage	7
9 Bonus Features	7
9.1 Spring Boot Actuator	8
9.2 Container Image	8
10 Conclusions	8
10.1 Lessons Learned	8
11 References	8

1 Objective

This laboratory focuses on designing and implementing a production-quality REST API using Java 21 and Spring Boot 3.3.x, following layered architecture principles. The *Blueprints* system manages geometric drawings identified by author and name, each composed of a list of 2D points.

The specific goals include:

- **Layered Architecture:** Model, Persistence, Service, and Controller separation.
- **PostgreSQL Migration:** Replace in-memory storage with a real relational database via JPA/Hibernate.
- **REST Best Practices:** Versioned paths (`/api/v1/`), correct HTTP codes, and uniform `ApiResponse<T>` wrappers.
- **OpenAPI/Swagger:** Auto-generated, annotated interactive documentation.
- **Configurable Filters:** Strategy pattern via Spring profiles (redundancy, undersampling).
- **Bonus:** Container image via `spring-boot:build-image` and production metrics with Spring Actuator.

2 Application Architecture

2.1 Package Structure

The project follows a strict **logical layers** pattern, ensuring that each layer depends only on the one below it:

Package	Responsibility
<code>model</code>	Domain entities: Blueprint, Point
<code>persistence</code>	Interface + InMemory/Postgres impls
<code>persistence.entity</code>	JPA entities: BlueprintEntity, PointEntity
<code>persistence.jpa</code>	Spring Data JPA repository
<code>services</code>	Business logic, filter orchestration
<code>filters</code>	Strategy: Identity, Redundancy, Undersampling
<code>controllers</code>	REST endpoints, DTO, ApiResponse
<code>config</code>	OpenAPI/Swagger configuration

Table 1: Project package structure

2.2 Request / Response Flow

Every API call traverses the following pipeline:

1. **Controller** validates the HTTP request and maps it to a service call.
2. **Service** applies the active filter and delegates to persistence.

3. **Persistence** reads or writes the data store (in-memory or PostgreSQL).
4. The result is wrapped in `ApiResponse<T>` and returned with the appropriate HTTP status.

Key Design Decision

Filters are applied on **all read operations** (`getAll`, `getByAuthor`, `getByName`). Write operations (`save`, `addPoint`) store raw data without filtering.

3 Part 1: Code Base Familiarization

3.1 Model Layer

- **Blueprint:** Identified uniquely by `author + name`. Exposes an unmodifiable view of its point list. Equality and `hashCode` are based only on author and name.
- **Point:** An immutable Java record holding `(int x, int y)`.

Listing 1: Immutable Point record

```
1 public record Point(
2     @Schema(description = "X coordinate") int x,
3     @Schema(description = "Y coordinate") int y
4 ) {}
```

3.2 Persistence Layer

`BlueprintPersistence` defines the storage contract:

Listing 2: Persistence interface contract

```
1 public interface BlueprintPersistence {
2     void saveBlueprint(Blueprint bp)
3         throws BlueprintPersistenceException;
4     Blueprint getBlueprint(String author, String name)
5         throws BlueprintNotFoundException;
6     Set<Blueprint> getBlueprintsByAuthor(String author)
7         throws BlueprintNotFoundException;
8     Set<Blueprint> getAllBlueprints();
9     void addPoint(String author, String name,
10         int x, int y) throws BlueprintNotFoundException;
11 }
```

`InMemoryBlueprintPersistence` implements this interface using a `ConcurrentHashMap` with composite key "`author:name`". Pre-loaded sample data:

Author	Name	Points
john	house	4
john	garage	3
jane	garden	3

Table 2: Sample data loaded on startup

3.3 Service Layer

`BlueprintsServices` acts as the orchestration layer. It injects both `BlueprintPersistence` and `BlueprintsFilter` via constructor injection, ensuring loose coupling and testability.

3.4 Controller Layer

BlueprintsAPIController maps all HTTP endpoints:

Method	Path	Code
GET	/api/v1/blueprints	200
GET	.../{author}	200/404
GET	.../{author}/{name}	200/404
POST	/api/v1/blueprints	201/400/403
PUT	.../{author}/{name}/points	202/404

Table 3: Controller endpoints and HTTP codes

4 Part 2: Migration to PostgreSQL

4.1 Design Decisions

The migration was implemented without modifying any existing class. Two key Spring annotations drive the behavior:

- `@Profile("postgres")`: The new implementation only loads when this profile is active.
- `@Primary`: When the profile is active, Spring selects `PostgresBlueprintPersistence` over `InMemoryBlueprintPersistence`.

Zero-Impact Migration

The full layered stack (Service, Controller, Filters) required **zero changes**. Only new classes were added in the persistence layer, honoring the Open/Closed Principle.

4.2 JPA Entity Model

Two JPA entities map to the relational schema:

Listing 3: BlueprintEntity – bidirectional mapping

```

1  @Entity
2  @Table(name = "blueprints",
3      uniqueConstraints = @UniqueConstraint(
4          columnNames = {"author", "name"}))
5  public class BlueprintEntity {
6      @OneToMany(mappedBy = "blueprint",
7          cascade = CascadeType.ALL,
8          orphanRemoval = true,
9          fetch = FetchType.EAGER)
10     @OrderColumn(name = "position")
11     private List<PointEntity> points = new ArrayList<>();
12
13 }
```

The `@OrderColumn` preserves point insertion order in the database, which is critical for filters like `UndersamplingFilter` that rely on positional indices.

4.3 Auto-Generated Relational Model

Table: blueprints

id	BIGINT	PK, auto-increment
author	VARCHAR	NOT NULL
name	VARCHAR	NOT NULL
UNIQUE(author, name)		

Table: points

id	BIGINT	PK, auto-increment
x	INT	NOT NULL
y	INT	NOT NULL
position	INT	Order in blueprint
blueprint_id	BIGINT	FK → blueprints(id)

Table 4: Auto-generated schema by Hibernate

4.4 Domain ↔ Entity Mapping

Listing 4: toDomain – entity to domain model

```

1  private Blueprint toDomain(BlueprintEntity e) {
2      List<Point> pts = e.getPoints().stream()
3          .map(p -> new Point(p.getX(), p.getY()))
4          .collect(Collectors.toList());
5      return new Blueprint(e.getAuthor(), e.getName(), pts);
6  }
```

Listing 5: toEntity – domain to JPA entity

```

1  private BlueprintEntity toEntity(Blueprint bp) {
2      BlueprintEntity e = new BlueprintEntity(
3          bp.getAuthor(), bp.getName());
4      bp.getPoints().forEach(p ->
5          e.addPoint(new PointEntity(p.x(), p.y())));
6      return e;
7  }
```

4.5 Docker Setup

Listing 6: Starting PostgreSQL with Docker

```

1  docker run --name blueprints-db \
2      -e POSTGRES_USER=postgres \
3      -e POSTGRES_PASSWORD=postgres \
4      -e POSTGRES_DB=blueprints \
5      -p 5432:5432 -d postgres:16
```

Important Note

On Windows, port 5432 may already be in use by a local PostgreSQL installation. If so, use `-p 5433:5432` and update the datasource URL in `application-postgres.properties` accordingly.

4.6 Database Verification

After running with `-Dspring-boot.run.profiles=postgres`, Hibernate auto-creates the schema. The data can be verified directly:

Listing 7: Verifying persisted blueprints

```

1  SELECT b.author, b.name, p.x, p.y, p.position
2  FROM blueprints b
3  JOIN points p ON p.blueprint_id = b.id
4  ORDER BY b.author, b.name, p.position;
```

```
blueprints=# SELECT * FROM blueprints;
 id | author | name
----+-----+
 1 | john   | house
 2 | john   | garage
 3 | jane   | garden
(3 rows)
```

Figure 1: PostgreSQL query confirming persisted blueprints

```
blueprints=# SELECT b.author, b.name, p.x, p.y, p.position
FROM blueprints b
JOIN points p ON p.blueprint_id = b.id
WHERE b.author = 'john' AND b.name = 'house'
ORDER BY p.position;
 author | name | x | y | position
-----+-----+---+---+-----
john  | house | 0 | 0 | 0
john  | house | 10 | 0 | 1
john  | house | 10 | 10 | 2
john  | house | 0 | 10 | 3
(4 rows)
```

Figure 2: Blueprint and its points joined from both tables

5 Part 3: REST API Best Practices

5.1 Versioned Base Path

All endpoints were moved from `/blueprints` to `/api/v1/blueprints`, enabling forward compatibility when the API evolves:

Listing 8: Versioned controller mapping

```
1 @RestController
2 @RequestMapping("/api/v1/blueprints")
3 public class BlueprintsAPIController { ... }
```

5.2 Uniform Response Wrapper

Every endpoint returns an `ApiResponse<T>`, a generic record that standardizes the response shape across the entire API:

Listing 9: ApiResponse record

```
1 public record ApiResponse<T>(
2     int code,
3     String message,
4     T data
5 ) {}
```

Example response for GET /api/v1/blueprints/john/house:

Listing 10: Successful 200 response

```
1 {
2     "code": 200,
3     "message": "execute ok",
4     "data": {
5         "author": "john",
6         "name": "house",
7         "points": [
8             {"x":0,"y":0}, {"x":10,"y":0},
```

```
9         {"x":10,"y":10}, {"x":0,"y":10}
10    ]
11 }
12 }
```

5.3 HTTP Status Code Mapping

Code	Scenario
200	Successful read
201	Blueprint created
202	Point added
400	Validation failure (missing fields)
403	Duplicate blueprint
404	Author or blueprint not found

Table 5: HTTP codes used by the API

5.4 Error Response Format

Errors also follow the `ApiResponse<T>` structure with `data: null`:

Listing 11: 404 error response

```
1 {
2     "code": 404,
3     "message": "Blueprint not found: john/xyz",
4     "data": null
5 }
```

6 Part 4: OpenAPI / Swagger

6.1 Configuration

The `springdoc-openapi-starter-webmvc-ui` dependency (v2.6.0) was added to `pom.xml`. `OpenApiConfig.java` customizes the metadata:

Listing 12: OpenAPI bean configuration

```
1 @Bean
2 public OpenAPI api() {
3     return new OpenAPI()
4         .info(new Info()
5             .title("ARSW Blueprints API")
6             .version("v1")
7             .description("REST API for blueprint management"))
8         .servers(List.of(new Server()
9             .url("http://localhost:8080")
10            .description("Development Server"))));
11 }
```

6.2 Endpoint Annotations

All controller methods are annotated with:

- `@Operation`: Summary and description.
- `@ApiResponses`: All possible response codes.
- `@Parameter`: Path variable descriptions and examples.
- `@Schema`: Request/response body examples.

6.3 Swagger UI Evidence

The screenshot shows the Swagger UI interface for the ARSW Blueprints API. It lists several endpoints under the 'Blueprints' category:

- GET /api/v1/blueprints**: Get all blueprints.
- GET /api/v1/blueprints/{author}**: Get blueprints by author.
- GET /api/v1/blueprints/{author}/{bpname}**: Get a specific blueprint.
- POST /api/v1/blueprints**: Create a new blueprint.
- PUT /api/v1/blueprints/{author}/{bpname}/points**: Add a point to a blueprint.

Figure 3: Swagger UI showing all versioned endpoints

The screenshot shows a successful GET request to the '/api/v1/blueprints' endpoint. The response body is a JSON array of blueprints, each with fields like 'id', 'name', 'author', 'x', 'y', and 'z'. The response code is 200 OK.

Figure 4: GET /api/v1/blueprints – successful response in Swagger

The screenshot shows a successful POST request to the '/api/v1/blueprints' endpoint. The response body is a JSON object representing a newly created blueprint with fields like 'id', 'name', 'author', 'x', 'y', and 'z'. The response code is 201 Created.

Figure 5: POST – 201 Created response with blueprint data

The screenshot shows a 404 Not Found response for a non-existent blueprint. The response body is a JSON object with fields 'code', 'message', 'status', and 'data'. The response code is 404.

Figure 6: GET non-existent blueprint – 404 Not Found

7 Part 5: Blueprints Filters

7.1 Design Pattern

Filters follow the **Strategy pattern**: all implement `BlueprintsFilter`, and Spring injects the active one based on the active profile. This avoids any `if/switch` logic in the service layer.

Listing 13: Filter interface

```
1 public interface BlueprintsFilter {
2     Blueprint apply(Blueprint bp);
3 }
```

7.2 Available Filters

Filter	Profile	Behavior
IdentityFilter	(default)	Returns blueprint unchanged
RedundancyFilter	redundancy	Removes consecutive duplicates
UndersamplingFilter	undersampling	Keeps even-indexed points

Table 6: Available filters and their activation profiles

7.3 RedundancyFilter

Iterates the point list sequentially and discards any point equal to its predecessor:

Listing 14: RedundancyFilter core logic

```
1 Point prev = null;
2 for (Point p : in) {
3     if (prev == null ||
4         !(prev.x() == p.x() &amp; prev.y() == p.y())) {
5         out.add(p);
6         prev = p;
7     }
8 }
```

Example: $(0,0), (0,0), (1,1), (1,1), (2,2) \rightarrow (0,0), (1,1), (2,2)$

7.4 UndersamplingFilter

Keeps only points at even indices $(0, 2, 4, \dots)$. Blueprints with ≤ 2 points are returned unchanged:

Listing 15: UndersamplingFilter core logic

```
1 for (int i = 0; i < in.size(); i++) {
2     if (i % 2 == 0) out.add(in.get(i));
3 }
```

Example: $(0,0), (1,1), (2,2), (3,3), (4,4) \rightarrow (0,0), (2,2), (4,4)$

7.5 Profile-Based Activation

The `IdentityFilter` uses a negated profile expression to avoid bean conflicts:

Listing 16: IdentityFilter profile guard

```
1 @Component
2 @Profile("!redundancy & !undersampling")
3 public class IdentityFilter implements BlueprintsFilter {
4     @Override
5     public Blueprint apply(Blueprint bp) { return bp; }
6 }
```

7.6 Filter Evidence




Figure 7: Identity filter: all 4 original points returned




Figure 8: Redundancy filter: consecutive duplicate points removed



Figure 9: Undersampling filter: 4 points reduced to 2 (even indices)

8 Testing Strategy

8.1 Test Suite Overview

Test Class	Layer	#
BlueprintsSmokeTest	Context	1
BlueprintModelTests	Model	7
BlueprintPersistenceTests	InMemory	8
BlueprintServiceTests	Service	9
FiltersTest	Filters	5
BlueprintControllerTests	Controller	9
OpenApiConfigTest	Config	2
BlueprintEntityTest	JPA Entity	4
PointEntityTest	JPA Entity	4
PostgresBlueprintPersistenceTest	Postgres	15
Total		64

Table 7: Test suite breakdown

8.2 Notable Test Techniques

MockMvc for Controller Tests – Loads the full Spring context without a real HTTP server, verifying JSON paths and HTTP status codes:

Listing 17: MockMvc controller test example

```
1  @Test
2  void testGetBlueprint() throws Exception {
3      mockMvc.perform(get("/api/v1/blueprints/john/house"))
4          .andExpect(status().isOk())
5          .andExpect(jsonPath("$.code").value(200))
6          .andExpect(jsonPath("$.data.author").value("john"));
7  }
```

Mockito for PostgresPersistence – Uses `@Mock` on `BlueprintJpaRepository` to test the full mapping logic without a database:

Listing 18: ArgumentCaptor verifies toEntity mapping

```
1  var captor = ArgumentCaptor.forClass(
2      BlueprintEntity.class);
3  verify(repo).save(captor.capture());
4  BlueprintEntity saved = captor.getValue();
5  assertEquals(10, saved.getPoints().get(0).getX());
```

8.3 SonarCloud Coverage

JaCoCo was configured without exclusions so that SonarCloud receives coverage data for all layers including JPA entities and the Postgres implementation. Key covered areas:

- All 5 BlueprintPersistence methods in Postgres impl.
- `toDomain` and `toEntity` private mapping helpers.
- All getters/setters in `BlueprintEntity` and `PointEntity`.
- `OpenApiConfig.api()` bean creation and server configuration.

9 Bonus Features

9.1 Spring Boot Actuator

Production-ready metrics and health endpoints were enabled by adding `spring-boot-starter-actuator` and configuring:

Listing 19: Actuator properties

```

1 management.endpoints.web.exposure.include=
2   health,info,metrics,env
3 management.endpoint.health.show-details=always
4 info.app.name=ARSW Blueprints API
5 info.app.version=1.0.0

```

Available endpoints at runtime:

Endpoint	URL
Health	/actuator/health
Info	/actuator/info
Metrics	/actuator/metrics
Env	/actuator/env

Table 8: Exposed Actuator endpoints

9.2 Container Image

Spring Boot's built-in Cloud Native Buildpacks support was used to generate an OCI image without a manual Dockerfile:

Listing 20: Building the container image

```
1 mvn spring-boot:build-image -DskipTests
```

The plugin was configured in `pom.xml`:

Listing 21: Image name configuration

```

1 <plugin>
2   <groupId>org.springframework.boot</groupId>
3   <artifactId>spring-boot-maven-plugin</artifactId>
4   <configuration>
5     <image>
6       <name>arsw-blueprints-api:${project.version}</name>
7     </image>
8   </configuration>
9 </plugin>

```

Running the container with a specific profile:

Listing 22: Running with undersampling filter

```

1 docker run -p 8080:8080 \
2   -e SPRING_PROFILES_ACTIVE=undersampling \
3   arsw-blueprints-api:1.0.0

```

Aspect	Dockerfile	build-image
Requires Dockerfile	Yes	No
Layered image	Manual	Automatic
JVM optimization	Manual	Auto-configured

Table 9: Dockerfile vs `spring-boot:build-image`

10 Conclusions

Key Findings

- Layered Architecture Pays Off:** The persistence migration required zero changes to the service and controller layers, demonstrating the value of interface-based design.
- Spring Profiles for Environment Switching:** A single `@Profile` annotation switches the entire persistence stack cleanly, eliminating configuration conditionals in business code.
- Strategy Pattern for Extensibility:** Adding a new filter requires only a new class implementing `BlueprintsFilter`, with no changes to existing code.
- Uniform Responses Improve API Consistency:** `ApiResponse<T>` ensures every endpoint — success or error — shares the same JSON contract, simplifying client-side handling.
- OpenAPI as Living Documentation:** Annotating endpoints with `@Operation` and `@ApiResponse` generates documentation that is always in sync with the code.
- Testability by Design:** Constructor injection and interface-based persistence allowed Mockito to fully test the Postgres implementation without a real database.

10.1 Lessons Learned

- `@Profile("!a & !b")` is the correct way to make a default bean that yields to alternatives.
- `@OrderColumn` is essential when point order matters downstream.
- JaCoCo exclusions hide real coverage gaps from SonarCloud — prefer writing tests over adding exclusions.
- `spring.autoconfigure.exclude` must be cleared in the postgres profile or JPA won't start.

11 References

- Spring Boot Reference Documentation (v3.3.x). <https://docs.spring.io/spring-boot/docs/3.3.x/reference/html/>
- Springdoc OpenAPI Documentation. <https://springdoc.org/>
- JaCoCo Java Code Coverage Library. <https://www.jacoco.org/jacoco/trunk/doc/>
- SonarCloud Documentation. <https://docs.sonarcloud.io/>
- Martin, R.C. (2017). *Clean Architecture*. Prentice Hall.
- Fowler, M. (2002). *Patterns of Enterprise Application Architecture*. Addison-Wesley.